

ANSHUMAN KUMAR

Bangalore · anshumankumar.mail@gmail.com · 9818164505 · anshumankumar.dev

WORK EXPERIENCE

Flexera, Bangalore

December 2025 – Present

Software Engineer III, FinOps AI/ML Enablement

- Own cloud-cost anomaly detection end to end: built it from scratch, replacing manual trial-and-error investigation with an automated daily scan of every enterprise client's full cloud spend, ranked by dollar impact and shipped with a plain-English explanation of what changed, by how much, and in which service, region, and account. Built on the FinOps Foundation's FOCUS specification, so detection is provider-agnostic across AWS, Azure, and GCP.
- Exposed that product to the company's customer-facing conversational AI assistant: built the MCP server and tool surface letting users view, acknowledge, and reject anomalies in natural language, over a Cube semantic layer on cost data. Also built a reusable GitHub Actions workflow that scans PRs and is now adopted by other engineers on the team.
- Own the root-cause-analysis agent that runs after an anomaly is flagged: independently re-investigates the anomaly window against raw cost data, scores every candidate service as a suspect, and returns either a specific top suspect or an honest "broad, org-wide shift" verdict when no single service stands out — so practitioners get a ranked, explained answer instead of a bare alert.
- Contribute to the architecture of the assistant as it moves to a multi-agent design: tool boundaries and ownership across agents, tracing and observability, response caching, and per-agent model selection.
- Cut projected pipeline operating cost by ~60% and bounded it per client — moved LLM enrichment off the detection critical path (inference wall-clock 1–2+ hours to ~20 minutes) and capped enrichment to the top-30 dollar-ranked anomalies per client — so cost stays flat as the fleet scales toward tens of thousands of clients. Authored the operating-cost analysis presented to senior leadership for the GA decision.
- Designed the pipeline to absorb extreme client skew and heterogeneous spend patterns: the largest enterprise client (3,000+ child billing organisations) runs as an isolated, independently-sized job so no single client inflates runtime for the fleet, while per-client model selection (Prophet, ARIMA, exponential smoothing) is automated across every meaningful combination of FOCUS cost dimensions via Spark GROUPING SETS, with no manual tuning. Deployed across three AWS production regions on Terraform-managed infrastructure.

Deloitte, Bangalore

December 2023 – December 2025

Senior Consultant, AI

- Architected a GenAI conversational-analytics agent with RAG and text-to-SQL for an OEM manufacturer (Mistral 8x7B, Codestral 22B, LLaVA-NeXT 34B; PGVector), answering questions over images, PDFs, and PLM manuals. Scaled the text-to-SQL layer from 3 to 100+ production tables — cut document retrieval time ~30% and improved query response times 20–25%.
- Built a supply-chain business-rule simulation and ML models predicting shipment cycle times and late/on-time status at 95% accuracy, with Feast for feature management at inference.
- Deployed AI services on bare-metal Kubernetes with GitLab CI; built an Azure OpenAI + Azure AI Search RAG chatbot giving field fibre mechanics step-by-step guidance.

Quantiphi, Bangalore

November 2020 – December 2023

Senior SDE

- Built chat/voice bots and KPI dashboards for government and private-sector clients (including a COVID-19 unemployment-insurance chatbot); early hands-on exposure to GenAI via PaLM on GCP/Vertex AI.

OYO Rooms

January 2020 – May 2020

Software Engineering Intern

SKILLS

- **Languages:** Python, SQL, JavaScript
- **Agentic AI & LLM:** MCP (server design, tool contracts), tool calling, multi-agent architecture, RAG, text-to-SQL, semantic layers (Cube), context design, AWS Bedrock, Vertex AI
- **ML:** Time-series forecasting (Prophet, ARIMA), anomaly detection, scikit-learn, Keras, MLflow, Feast
- **Backend & Cloud:** Docker, Kubernetes, Terraform, AWS, GCP, GitLab CI, GitHub Actions, Databricks, PySpark, Delta Lake, FOCUS specification
- **Databases:** Postgres, MySQL, Firestore, DocumentDB
- **Certifications:** AWS Certified Developer – Associate (2024–2027); GCP Professional Cloud Developer (2022–2024)

EDUCATION

- **Manipal Institute of Technology, Manipal** — B.Tech, Computer Science and Engineering — GPA 7.35 2016 – 2020
- **The Mother's International School, New Delhi** — CBSE — 93.6%

ACHIEVEMENTS & EXTRA-CURRICULARS

- Built an open-source coding-automation tool (unit-test generation, code fixes, Stack Overflow search, code completion) — third prize, 24-hour Jovian Hackathon 2023, Bangalore.
- Conceived and shipped Keystone, a published Chrome extension that uses AI to assess webpages against a user's stated focus goal and deliver adaptive nudges during timed focus sessions — built end-to-end using AI-assisted development outside a primarily backend/Python day job.
- Built an MCP server (Python, AWS Lambda) giving AI assistants CMS-style publishing tools — create, update, and delete posts on a Git-backed static blog — as a lightweight, git-native alternative to a hosted CMS.
- Cycling and marathons — Vedanta Delhi Half Marathon (2022), Bangalore Half Marathon (2023), and a 200 km cycling event.